

There are some queues we have access to that you didn't list in the doc (all the new-* ones).

The new-* are queues without GPU?

I think, yes. I checked which hosts are inside new* queues:

blimits -a -q waic-short	limits slots					
	limit ngpus_phys					
		new-risk	new-long	new-medium	new-interactive	new-short
time limit	min	100800	10080	4320	1440	1440
	h	1680	168	72	24	24
	days	70	7	3	1	1
	hosts_group	shalev_hosts/	shalev_hosts/	shalev_hosts/		shalev_hosts/
		elinav_hosts/	elinav_hosts/	elinav_hosts/		elinav_hosts/
bqueues -l new-risk		ginossar_hosts/	ginossar_hosts/	ginossar_hosts/		ginossar_hosts/
		public_himem_hosts/	public_himem_hosts/	public_himem_hosts/		public_himem_hosts/
		dell_hosts/	dell_hosts/	dell_hosts/		dell_hosts/
		schwartz_hosts/	schwartz_hosts/	schwartz_hosts/		schwartz_hosts/
		shlush_hosts/	shlush_hosts/	shlush_hosts/		shlush_hosts/
		bio_hosts/	bio_hosts/	bio_hosts/		bio_hosts/
		wicc_hosts/	wicc_hosts/	wicc_hosts/		wicc_hosts/
		ulitsky_hosts/	ulitsky_hosts/	ulitsky_hosts/		ulitsky_hosts/
		molgen_2020_hosts/	molgen_2020_hosts/	molgen_2020_hosts/		molgen_2020_hosts/
		ulitsky_2020_hosts/	ulitsky_2020_hosts/	ulitsky_2020_hosts/		ulitsky_2020_hosts/
		fleishman_2020_hosts/	fleishman_2020_hosts/	fleishman_2020_hosts/		fleishman_2020_hosts/
		feinerman_2020_hosts/	feinerman_2020_hosts/	feinerman_2020_hosts/		feinerman_2020_hosts/
		public_2017_hosts/	public_2017_hosts/	public_2017_hosts/		public_2017_hosts/
		isilon_hosts/	isilon_hosts/	isilon_hosts/		isilon_hosts/
		yan_hosts/	yan_hosts/	yan_hosts/		yan_hosts/
		samuels_2021_hosts/	samuels_2021_hosts/	samuels_2021_hosts/		samuels_2021_hosts/
		public_hosts/	public_hosts/	public_hosts/	public_hosts	public_hosts/
		berg_hosts/	berg_hosts/	berg_hosts/		berg_hosts/
		schneidman_hosts/	schneidman_hosts/	schneidman_hosts/		schneidman_hosts/
		merbl_hosts/	merbl_hosts/	merbl_hosts/		merbl_hosts/
		physics_2021_hosts/	physics_2021_hosts/	physics_2021_hosts/		physics_2021_hosts/
		bio_hosts_2022/	bio_hosts_2022/	bio_hosts_2022/		bio_hosts_2022/
		tirosch_hosts/	tirosch_hosts/	tirosch_hosts/		tirosch_hosts/
						cn239
		public_gsla_hosts/	public_gsla_hosts/	public_gsla_hosts/		public_gsla_hosts/
		public_gsla_mem_hosts/	public_gsla_mem_hosts/	public_gsla_mem_hosts/		public_gsla_mem_hosts/
		koren_2020_hosts/	koren_2020_hosts/	koren_2020_hosts/		koren_2020_hosts/
		CPU	CPU	CPU	CPU	CPU

and checked, for example the host group = ginossar_host => it includes two hosts: cn242 and cn243:

```
[ingap@access4 scripts]$ bhosts ginossar_hosts
```

HOST NAME	STATUS	JL/U	MAX	NJOBS	RUN	SSUSP	USUSP	RSV
cn242	ok	-	104	21	21	0	0	0
cn243	ok	-	104	49	49	0	0	0

I run the command: ssh cn242 'nvidia-smi' => to see which type of GPU is there => the output was:

```
[ingap@access4 scripts]$ ssh cn242 'nvidia-smi'
```

```
Warning: Permanently added 'cn242,132.76.221.17' (ECDSA) to the list of known hosts.
```

```
ingap@cn242's password:
```

```
bash: nvidia-smi: command not found
```

or you can run:

```
[ingap@access4 scripts]$ lshosts -gpu cn242
```

```
lshosts: Waiting for LIM to collect GPU resource usage values or no hosts with GPU devices are detected.
```

so, my conclusion => it is not GPU.

There are additional queue which includes GPU, like yan-gpu with even higher priority than waic* queue:

```
[ingap@access4 ~]$ bqueues
```

QUEUE_NAME	PRIO	STATUS	MAX	JL/U	JL/P	JL/H	NJOBS	PEND	RUN	SUSP
molgen-short	302	Open:Active	-	50	-	-	0	0	0	0
molgen-g	300	Open:Active	-	700	-	-	40	0	40	0
yan-gpu	289	Open:Active	-	-	-	-	0	0	0	0
sch-gpu	283	Open:Active	-	-	-	-	1	0	1	0
gsia-mem	240	Open:Active	-	-	-	-	0	0	0	0
gsia-cpu	240	Open:Active	-	-	-	-	5617	3744	1873	0
high_gpu_gsla	240	Open:Active	-	10	-	-	1	0	1	0
interactive_cpu	240	Open:Active	-	-	-	-	0	0	0	0
interactive_gpu	240	Open:Active	-	10	-	-	0	0	0	0
waic-short	189	Open:Active	-	-	-	-	7	0	7	0
waic-medium	188	Open:Active	-	-	-	-	2	0	2	0
waic-long	186	Open:Active	-	-	-	-	37	12	25	0
waic-risk	185	Open:Active	-	-	-	-	267	54	213	0
elinav	180	Open:Active	-	-	-	-	9	9	0	0
sorek-gpu	145	Open:Active	-	-	-	-	0	0	0	0
leeat-gpu	145	Open:Active	-	-	-	-	0	0	0	0
molgen-gpu	140	Open:Active	-	-	-	-	1	0	1	0
bio-gpu	140	Open:Active	-	-	-	-	0	0	0	0
gsia-gpu	140	Open:Active	-	-	-	-	1	0	1	0
gpu-test	90	Open:Active	-	-	-	-	0	0	0	0
gpu-interactive	89	Open:Active	-	30	-	-	6	0	6	0
gpu-shared	85	Closed:Active	-	1000	-	-	0	0	0	0
gpu-short	84	Open:Active	-	1000	-	-	69	0	69	0
gpu-medium	84	Open:Active	-	720	-	-	23	0	23	0
gpu-long	84	Open:Active	-	460	-	-	478	438	40	0
gpu-risk	83	Open:Active	-	1000	-	-	0	0	0	0
gpu-parallel	82	Open:Active	-	1000	-	-	0	0	0	0

yan-gpu queue includes 2 hosts ibdgx011 and ibdgx012:

```
[ingap@access4 scripts]$ bqueues -l yan-gpu

QUEUE: yan-gpu
-- queue for Yan Binghai GPU jobs

PARAMETERS/STATISTICS
PRIO NICE STATUS      MAX JL/U JL/P JL/H NJOBS  PEND   RUN  SSUSP  USUSP  RSV  PJOBS
289  20  Open:Active        -  -  -  -  -  0      0      0      0      0      0
Interval for a host to accept two jobs is 0 seconds
RUNLIMIT
1440.0 min

SCHEDULING PARAMETERS
r15s  r1m  r15m  ut    pg    io    ls    it    tmp    swp    mem
loadSched - - - - - - - - - - -
loadStop - - - - - - - - - - -

SCHEDULING POLICIES: FAIRSHARE
USER_SHARES: [default, 1]

HIST_HOURS: 72.000000

USERS: talma yan-wx-grp/ wexac-admins/
HOSTS: ibdgx011 ibdgx012
RES_REQ: select[type=any] order[-mem]
RESRSV_LIMIT: [ngpus_physical=0,10]
GPU_REQ: num=1
ENABLE_GPU_HIST_RUN_TIME: Y

[ingap@access4 scripts]$
```

and these hosts are GPUs =>

```
[ingap@access4 scripts]$ lshosts -gpu ibdgx011
```

HOST_NAME	gpu_id	gpu_model	gpu_driver	gpu_factor	numa_id	vendor	mig
ibdgx011	0	TeslaV100_SXM2	525.60.13	7.0	0	Nvidia	N
	1	TeslaV100_SXM2	525.60.13	7.0	0	Nvidia	N
	2	TeslaV100_SXM2	525.60.13	7.0	0	Nvidia	N
	3	TeslaV100_SXM2	525.60.13	7.0	0	Nvidia	N
	4	TeslaV100_SXM2	525.60.13	7.0	1	Nvidia	N
	5	TeslaV100_SXM2	525.60.13	7.0	1	Nvidia	N
	6	TeslaV100_SXM2	525.60.13	7.0	1	Nvidia	N
	7	TeslaV100_SXM2	525.60.13	7.0	1	Nvidia	N

However, we can not use this queue:

```
[ingap@access4 scripts]$ bqueues -l yan-gpu

QUEUE: yan-gpu
-- queue for Yan Binghai GPU jobs

PARAMETERS/STATISTICS
PRIO NICE STATUS          MAX JL/U JL/P JL/H NJOBS  PEND   RUN  SSUSP  USUSP   RSV  PJOBS
289  20  Open:Active          -   -   -   -   0      0     0     0     0     0     0
Interval for a host to accept two jobs is 0 seconds
RUNLIMIT
1440.0 min

SCHEDULING PARAMETERS
      r15s  r1m  r15m  ut      pg   io   ls   it   tmp   swp   mem
loadSched -   -   -   -      -   -   -   -   -   -   -
loadStop  -   -   -   -      -   -   -   -   -   -   -

SCHEDULING POLICIES: FAIRSHARE
USER_SHARES: [default, 1]

HIST_HOURS: 72.000000
USERS: talma yan-wx-grp/ wexac-admins/
HOSTS: ibdgx011 ibdgx012
RES_REQ: select[type=any].order[-mem]
RESRSV_LIMIT: [ngpus_physical=0,10]
GPU_REQ: num=1
ENABLE_GPU_HIST_RUN_TIME: Y

[ingap@access4 scripts]$
```

so, our GPUs are waic* and gpu*.

I forgot to mention in the Task 3 that waic-risk can be preempted by waic-short/long/medium

Can you explain how the "Factor" works?

According to IBM LSF manual:

<https://www.ibm.com/docs/en/spectrum-lsf/10.1.0?topic=reference-command>

<https://www.ibm.com/docs/en/spectrum-lsf/10.1.0?topic=reference-bqueues>

RUN_TIME_FACTOR

The weighting parameter for run_time within the dynamic priority calculation. If not defined for the queue, the cluster-wide value that is defined in the `lsb.params` file is used.

The explanation is not clear for me. Don't know. Lets leave this at his stage.

Does "RESRSV_LIMIT: [ngpus_physical=0,8]" means we cant use less than 0 gpus or more than 8?

I think yes, this is range for amount of GPU which can be used per job: from 0 till 8.

Can you explain how the "nfs4_editfacl" command works? (the one from the jupyter tutorial in the hpcwiki)

I didn't have a time yet to read HPC wiki. I will send the task related to Jupyter how I open and work with it => how I do it.